

# Deep Learning for Perception

## Lecture 03: Stochastic Gradient Descent



### Incremental Learning

#### Topics Covered in This Lecture:

- Review: Batch Gradient Descent
- Stochastic Gradient Descent (SGD)
- SGD Algorithm Step-by-Step
- SGD vs Batch GD Comparison
- Advantages of SGD
- Challenges & Solutions
- Mini-Batch Gradient Descent
- Multiple Numerical Examples

**Instructor:** Aqsa Younas

Department of Computer Science  
FAST-NU, CFD Campus

## Contents

<b>1 Review: Batch Gradient Descent</b>	<b>2</b>
1.1 Batch GD Recap . . . . .	2
1.2 Problems with Batch GD . . . . .	3
<b>2 Stochastic Gradient Descent (SGD)</b>	<b>4</b>
2.1 Core Idea of SGD . . . . .	4
2.2 SGD Formula . . . . .	4
2.3 SGD Error Function . . . . .	5
<b>3 SGD Algorithm: Step-by-Step</b>	<b>6</b>
<b>4 Numerical Example 1: Basic SGD</b>	<b>8</b>
<b>5 SGD vs Batch Gradient Descent</b>	<b>11</b>
5.1 Visual Comparison of Convergence Paths . . . . .	11
5.2 Detailed Comparison Table . . . . .	11
5.3 Advantages of SGD . . . . .	12
<b>6 Numerical Example 2: SGD vs Batch GD Comparison</b>	<b>13</b>
<b>7 SGD Over Multiple Epochs</b>	<b>16</b>
<b>8 Mini-Batch Gradient Descent</b>	<b>19</b>
8.1 Mini-Batch GD Concept . . . . .	19
8.2 Mini-Batch GD Advantages . . . . .	19
<b>9 Practical Considerations for SGD</b>	<b>22</b>
9.1 Learning Rate Selection . . . . .	22
9.2 When SGD Can Escape Local Minima . . . . .	22
<b>10 Summary: SGD Calculation Checklist</b>	<b>23</b>
<b>11 Glossary</b>	<b>24</b>

### Advance Organizer — What You'll Learn

**Learning Objectives:** By the end of this lecture, you will be able to:

1. **Explain** the motivation behind Stochastic Gradient Descent
2. **Implement** SGD weight updates step-by-step
3. **Compare** SGD with Batch Gradient Descent mathematically
4. **Calculate** weight updates for multiple epochs
5. **Understand** Mini-Batch GD as a compromise
6. **Solve** numerical problems involving SGD from scratch

**Prior Knowledge Required:**

- Batch Gradient Descent algorithm
- Weight update rule:  $\Delta w_i = \eta(t - o)x_i$
- Loss functions (MSE)

## 1 Review: Batch Gradient Descent

### Why It Matters

Before understanding SGD, we must clearly understand **Batch Gradient Descent** and its limitations. SGD was developed to address specific problems with Batch GD.

### 1.1 Batch GD Recap

#### Definition

**Batch Gradient Descent** computes the gradient using **ALL** training samples before making a single weight update.

#### Process:

1. Process **every sample** in the dataset
2. **Accumulate** all weight changes ( $\Delta w$ )
3. Update weights **ONCE** at the end

### Key Formula

#### Batch GD Update Rule:

$$\Delta w_i = \eta \sum_{d \in D} (t_d - o_d) \cdot x_{id}$$

$$w_i \leftarrow w_i + \Delta w_i \quad (\text{update ONCE per epoch})$$

## 1.2 Problems with Batch GD

### Limitations of Batch Gradient Descent

1. **Computationally Expensive:** Must process ALL samples before one update
2. **Slow Convergence:** Only one weight update per epoch
3. **Memory Intensive:** Must store gradients for all samples
4. **Stuck in Local Minima:** Smooth path can get trapped
5. **Not Suitable for Large Datasets:** Impractical for millions of samples

## 2 Stochastic Gradient Descent (SGD)

### Why It Matters

SGD was developed to make learning **faster** and more **practical** for large datasets. Instead of waiting to see all data, we update weights **immediately** after each sample!

### 2.1 Core Idea of SGD

#### Definition

**Stochastic Gradient Descent (SGD)** updates weights **incrementally** after processing **each individual training example**, rather than waiting for the entire dataset.

**Key Insight:** Each sample provides an **estimate** of the true gradient. We use this noisy estimate to make progress immediately.

#### Analogy — Think of It Like This

##### Batch GD vs SGD — A Navigation Analogy:

**Batch GD:** Like planning a road trip by studying the **entire map** before taking your first step. You know the exact direction, but it takes forever to start moving.

**SGD:** Like asking **one local person** at a time for directions. Each direction might be slightly off, but you start moving immediately and eventually reach your destination through many small corrections.

### 2.2 SGD Formula

#### Key Formula

##### SGD Update Rule (for each sample $d$ ):

$$\Delta w_i = \eta(t_d - o_d) \cdot x_{id}$$

$$w_i \leftarrow w_i + \Delta w_i \quad (\text{update IMMEDIATELY})$$

Where:

- $\eta$  = Learning rate
- $t_d$  = Target value for sample  $d$
- $o_d$  = Output (prediction) for sample  $d$
- $x_{id}$  = Input feature  $i$  of sample  $d$
- $(t_d - o_d)$  = Error for sample  $d$

**Critical Difference:** Weights are updated **after EACH sample**, not after all samples!

## 2.3 SGD Error Function

### Definition

In SGD, we define a **per-sample error function**:

$$E_d(\mathbf{w}) = \frac{1}{2}(t_d - o_d)^2$$

Instead of minimizing total error  $E(\mathbf{w}) = \sum_d E_d(\mathbf{w})$ , SGD takes steps based on individual  $E_d(\mathbf{w})$  values.

### 3 SGD Algorithm: Step-by-Step

#### Stochastic Gradient Descent Algorithm

**Input:** Training data  $D$ , learning rate  $\eta$ , initial weights  $w$

**Repeat** (for each epoch):

1. **Shuffle** the training data (optional but recommended)
2. **For each** training sample  $(x, t)$  in  $D$ :
  - a. Calculate output:  $o = \sum_i w_i \cdot x_i$
  - b. Calculate error:  $e = t - o$
  - c. **For each** weight  $w_i$ :
    - Calculate delta:  $\Delta w_i = \eta \cdot e \cdot x_i$
    - **Update immediately**:  $w_i \leftarrow w_i + \Delta w_i$

**Until** convergence (error below threshold or max epochs reached)




Figure 1: SGD Algorithm Flowchart: Update weights after EACH sample

### Memory Hook — Remember This!

#### The Key Difference to Remember:

Batch GD	SGD
See ALL samples → Update ONCE	See ONE sample → Update ONCE
$n$ samples = 1 update	$n$ samples = $n$ updates

## 4 Numerical Example 1: Basic SGD

### Solved Example 1: SGD — One Epoch

#### Given Data:

- Learning rate:  $\eta = 0.5$
- Initial weights:  $w_0 = -0.3$  (bias),  $w_1 = 0.5$ ,  $w_2 = 0.5$
- Input format:  $x = [1, x_1, x_2]$  (first element is 1 for bias)

#### Training Data:

Sample	Input $x$	Target $t$
1	[1, 1, 0]	1
2	[1, 1, 1]	0
3	[1, 0, 0]	0
4	[1, 0, 1]	1

**Task:** Perform ONE epoch of SGD. Show all weight updates.

**Solution:****Output formula:**  $o = w_0 \cdot x_0 + w_1 \cdot x_1 + w_2 \cdot x_2$ **Update rule:**  $\Delta w_i = \eta \cdot (t - o) \cdot x_i$  (apply immediately!)**STEP 1: Process Sample 1 —  $x = [1, 1, 0]$ ,  $t = 1$** **Current weights:**  $[w_0, w_1, w_2] = [-0.3, 0.5, 0.5]$ **Forward Pass:**

$$o = (-0.3)(1) + (0.5)(1) + (0.5)(0) = -0.3 + 0.5 + 0 = \mathbf{0.2}$$

**Error:**

$$e = t - o = 1 - 0.2 = \mathbf{0.8}$$

**Weight Updates:**

$$\Delta w_0 = 0.5 \times 0.8 \times 1 = \mathbf{0.4}$$

$$w_0 = -0.3 + 0.4 = \mathbf{0.1}$$

$$\Delta w_1 = 0.5 \times 0.8 \times 1 = \mathbf{0.4}$$

$$w_1 = 0.5 + 0.4 = \mathbf{0.9}$$

$$\Delta w_2 = 0.5 \times 0.8 \times 0 = \mathbf{0}$$

$$w_2 = 0.5 + 0 = \mathbf{0.5}$$

Updated weights:  $[0.1, 0.9, 0.5]$ **STEP 2: Process Sample 2 —  $x = [1, 1, 1]$ ,  $t = 0$** **Current weights:**  $[0.1, 0.9, 0.5]$  (using UPDATED weights!)**Forward Pass:**

$$o = (0.1)(1) + (0.9)(1) + (0.5)(1) = 0.1 + 0.9 + 0.5 = \mathbf{1.5}$$

**Error:**

$$e = t - o = 0 - 1.5 = \mathbf{-1.5}$$

**Weight Updates:**

$$\Delta w_0 = 0.5 \times (-1.5) \times 1 = \mathbf{-0.75}$$

$$w_0 = 0.1 - 0.75 = \mathbf{-0.65}$$

$$\Delta w_1 = 0.5 \times (-1.5) \times 1 = \mathbf{-0.75}$$

$$w_1 = 0.9 - 0.75 = \mathbf{0.15}$$

$$\Delta w_2 = 0.5 \times (-1.5) \times 1 = \mathbf{-0.75}$$

$$w_2 = 0.5 - 0.75 = \mathbf{-0.25}$$

Updated weights:  $[-0.65, 0.15, -0.25]$ **STEP 3: Process Sample 3 —  $x = [1, 0, 0]$ ,  $t = 0$** **Current weights:**  $[-0.65, 0.15, -0.25]$ **Forward Pass:**

$$o = (-0.65)(1) + (0.15)(0) + (-0.25)(0) = \mathbf{-0.65}$$

**Error:**

$$e = t - o = 0 - (-0.65) = \mathbf{0.65}$$

**Weight Updates:**

$$\Delta w_0 = 0.5 \times 0.65 \times 1 = \mathbf{0.325}$$

$$w_0 = -0.65 + 0.325 = \mathbf{-0.325}$$

$$\Delta w_1 = 0.5 \times 0.65 \times 0 = \mathbf{0}$$

$$w_1 = 0.15 + 0 = \mathbf{0.15}$$

$$\Delta w_2 = 0.5 \times 0.65 \times 0 = \mathbf{0}$$

$$w_2 = -0.25 + 0 = \mathbf{-0.25}$$

FAST-NUCES — Stochastic Gradient Descent

Updated weights:  $[-0.325, 0.15, -0.25]$

**Final Answer (After 1 Epoch of SGD):**

$$w_0 = 0.4625, \quad w_1 = 0.15, \quad w_2 = 0.5375$$

**Key Observation:** Notice how each step used the **updated weights** from the previous step. This is what makes SGD different from Batch GD!




Figure 2: Weight evolution through SGD: 4 updates in 1 epoch

## 5 SGD vs Batch Gradient Descent

### Why It Matters

Understanding the differences between SGD and Batch GD helps you choose the right algorithm for your problem and understand why modern deep learning prefers SGD-based methods.

### 5.1 Visual Comparison of Convergence Paths




Figure 3: Convergence paths: Batch GD (smooth) vs SGD (noisy but often faster)

### 5.2 Detailed Comparison Table

Aspect	Batch GD	Stochastic GD
<b>Update Frequency</b>	Once per epoch	After every sample
<b>Updates per Epoch</b>	1 update	$n$ updates (where $n$ = samples)
<b>Gradient Used</b>	True gradient (exact)	Noisy estimate (approximate)
<b>Convergence Path</b>	Smooth, direct	Noisy, oscillating
<b>Speed (large data)</b>	Very slow	Much faster
<b>Memory Usage</b>	High (store all gradients)	Low (one sample at a time)
<b>Local Minima</b>	Can get stuck	Can escape (noise helps!)
<b>Learning Rate</b>	Can be larger	Usually needs smaller
<b>Best For</b>	Small datasets, convex problems	Large datasets, neural networks

Table 1: Comprehensive comparison of Batch GD vs SGD

### 5.3 Advantages of SGD

#### Why SGD is Preferred in Deep Learning

1. **Faster Progress:**  $n$  weight updates per epoch instead of 1
2. **Memory Efficient:** Only need to store one sample at a time
3. **Escapes Local Minima:** Noise in gradient can jump out of shallow valleys
4. **Online Learning:** Can learn from streaming data in real-time
5. **Regularization Effect:** Noise acts as implicit regularization

## 6 Numerical Example 2: SGD vs Batch GD Comparison

### Solved Example 2: Comparing SGD and Batch GD Results

#### Given Data (same for both):

- Learning rate:  $\eta = 0.5$
- Initial weights:  $w_0 = -0.3, w_1 = 0.5, w_2 = 0.5$
- Same 4 training samples as Example 1

**Task:** Compare final weights after 1 epoch using both methods.

**Method 1: Batch GD** (accumulate all, update once)Process all samples with **original weights**  $[-0.3, 0.5, 0.5]$ :

Sample	x	t	o	error	$\Delta w_0$	$\Delta w_1, \Delta w_2$
1	[1,1,0]	1	0.2	0.8	0.4	0.4, 0
2	[1,1,1]	0	0.7	-0.7	-0.35	-0.35, -0.35
3	[1,0,0]	0	-0.3	0.3	0.15	0, 0
4	[1,0,1]	1	0.2	0.8	0.4	0, 0.4
<b>Sum of <math>\Delta w</math>:</b>					<b>0.6</b>	<b>0.05, 0.05</b>

**Batch GD Final Weights:**

$$w_0 = -0.3 + 0.6 = \mathbf{0.3}$$

$$w_1 = 0.5 + 0.05 = \mathbf{0.55}$$

$$w_2 = 0.5 + 0.05 = \mathbf{0.55}$$

**Method 2: SGD** (from Example 1)Weights evolve:  $[-0.3, 0.5, 0.5] \rightarrow [0.1, 0.9, 0.5] \rightarrow [-0.65, 0.15, -0.25] \rightarrow [-0.325, 0.15, -0.25] \rightarrow [0.4625, 0.15, 0.5375]$ **SGD Final Weights:**  $w_0 = 0.4625, w_1 = 0.15, w_2 = 0.5375$ **Calculate MSE for Both:****Batch GD** with weights  $[0.3, 0.55, 0.55]$ :

$$\text{Sample 1: } o = 0.3 + 0.55 + 0 = 0.85, \quad e^2 = (1 - 0.85)^2 = 0.0225$$

$$\text{Sample 2: } o = 0.3 + 0.55 + 0.55 = 1.4, \quad e^2 = (0 - 1.4)^2 = 1.96$$

$$\text{Sample 3: } o = 0.3 + 0 + 0 = 0.3, \quad e^2 = (0 - 0.3)^2 = 0.09$$

$$\text{Sample 4: } o = 0.3 + 0 + 0.55 = 0.85, \quad e^2 = (1 - 0.85)^2 = 0.0225$$

$$\text{MSE} = \frac{0.0225 + 1.96 + 0.09 + 0.0225}{4} = \mathbf{0.5238}$$

**SGD** with weights  $[0.4625, 0.15, 0.5375]$ :

$$\text{Sample 1: } o = 0.4625 + 0.15 + 0 = 0.6125, \quad e^2 = (1 - 0.6125)^2 = 0.1502$$

$$\text{Sample 2: } o = 0.4625 + 0.15 + 0.5375 = 1.15, \quad e^2 = (0 - 1.15)^2 = 1.3225$$

$$\text{Sample 3: } o = 0.4625, \quad e^2 = (0 - 0.4625)^2 = 0.2139$$

$$\text{Sample 4: } o = 0.4625 + 0 + 0.5375 = 1.0, \quad e^2 = (1 - 1.0)^2 = 0$$

$$\text{MSE} = \frac{0.1502 + 1.3225 + 0.2139 + 0}{4} = \mathbf{0.4216}$$

**Comparison Results:**

Method	$w_0$	$w_1$	$w_2$	MSE
Batch GD	0.3	0.55	0.55	0.5238
SGD	0.4625	0.15	0.5375	<b>0.4216</b>

**Observation:** After just 1 epoch, SGD achieved **lower MSE** (0.4216 vs 0.5238). This is because SGD made 4 weight updates while Batch GD made only 1!

## 7 SGD Over Multiple Epochs

### Why It Matters

In practice, SGD runs for **multiple epochs**. The noisy updates gradually lead to convergence. Let's see how weights evolve over several passes through the data.

### Solved Example 3: SGD for 3 Epochs

**Given Data:** Same as Examples 1 and 2.

**Task:** Perform 3 complete epochs of SGD and track weight evolution.

**Solution:**

We continue from the same initial weights and show the summary for each epoch.

**EPOCH 1:** (Already computed in Example 1)

Step	Input	t	Error	Weights After
1	[1,1,0]	1	0.8	[0.1, 0.9, 0.5]
2	[1,1,1]	0	-1.5	[-0.65, 0.15, -0.25]
3	[1,0,0]	0	0.65	[-0.325, 0.15, -0.25]
4	[1,0,1]	1	1.575	[0.4625, 0.15, 0.5375]

**End of Epoch 1:** [0.4625, 0.15, 0.5375]

---

**EPOCH 2:** Starting with [0.4625, 0.15, 0.5375]

**Step 1:**  $x = [1, 1, 0], t = 1$

- $o = 0.4625 + 0.15 + 0 = 0.6125$
- Error =  $1 - 0.6125 = 0.3875$
- $\Delta w = [0.1938, 0.1938, 0]$
- Updated: [0.6563, 0.3438, 0.5375]

**Step 2:**  $x = [1, 1, 1], t = 0$

- $o = 0.6563 + 0.3438 + 0.5375 = 1.5375$
- Error =  $0 - 1.5375 = -1.5375$
- Updated: [-0.1125, -0.4250, -0.2313]

**Step 3:**  $x = [1, 0, 0], t = 0$

- $o = -0.1125$ , Error = 0.1125
- Updated: [-0.0563, -0.4250, -0.2313]

**Step 4:**  $x = [1, 0, 1], t = 1$

- $o = -0.0563 + 0 - 0.2313 = -0.2875$
- Error =  $1 - (-0.2875) = 1.2875$
- Updated: [0.5875, -0.4250, 0.4125]

**End of Epoch 2:** [0.5875, -0.4250, 0.4125]

---

**EPOCH 3:** Starting with [0.5875, -0.4250, 0.4125]

(Following same process...)

**End of Epoch 3:** [0.7219, -0.7125, 0.2781]

**Weight Evolution Summary:**

Epoch	$w_0$	$w_1$	$w_2$
0 (Initial)	-0.3	0.5	0.5
1	0.4625	0.15	0.5375
2	0.5875	-0.425	0.4125
3	0.7219	-0.7125	0.2781

**Observation:** Weights continue to change significantly even after multiple epochs. SGD's noisy path gradually refines the solution.

## 8 Mini-Batch Gradient Descent

### Why It Matters

Mini-Batch GD is a **compromise** between Batch GD and SGD. It offers the benefits of both: stable gradient estimates like Batch GD and frequent updates like SGD.

### 8.1 Mini-Batch GD Concept

#### Definition

**Mini-Batch Gradient Descent** divides the training data into small **batches** of size  $B$ . Weights are updated after processing each batch.

$$\Delta w_i = \eta \sum_{d \in \text{batch}} (t_d - o_d) \cdot x_{id}$$

**Common batch sizes:** 16, 32, 64, 128, 256

#### Batch GD



#### Mini-Batch (B=2)



#### SGD (B=1)



Figure 4: Comparison: Batch GD (1 update), Mini-Batch (4 updates), SGD (8 updates)

### 8.2 Mini-Batch GD Advantages

#### Why Mini-Batch is Often Best

- **Better gradient estimate** than SGD (averages over batch)
- **More updates** than Batch GD (faster progress)
- **GPU-friendly**: Batches can be processed in parallel
- **Memory efficient**: Only load batch into memory
- **Sweet spot**: Balances noise and stability

### Solved Example 4: Mini-Batch GD (Batch Size = 2)

#### Given Data:

- Learning rate:  $\eta = 0.5$ , Batch size:  $B = 2$
- Initial weights:  $w_0 = -0.3, w_1 = 0.5, w_2 = 0.5$
- Same 4 training samples

**Task:** Perform 1 epoch of Mini-Batch GD.

#### Solution:

Divide 4 samples into 2 mini-batches of size 2.

##### Mini-Batch 1: Samples 1 and 2

Using weights  $[-0.3, 0.5, 0.5]$ :

**Sample 1:**  $x = [1, 1, 0], t = 1$

- $o = -0.3 + 0.5 + 0 = 0.2$ , Error = 0.8
- $\Delta w = [0.4, 0.4, 0]$

**Sample 2:**  $x = [1, 1, 1], t = 0$

- $o = -0.3 + 0.5 + 0.5 = 0.7$ , Error = -0.7
- $\Delta w = [-0.35, -0.35, -0.35]$

**Accumulated  $\Delta w$ :**  $[0.4 - 0.35, 0.4 - 0.35, 0 - 0.35] = [0.05, 0.05, -0.35]$

**Update:**  $[-0.3 + 0.05, 0.5 + 0.05, 0.5 - 0.35] = [-0.25, 0.55, 0.15]$

##### Mini-Batch 2: Samples 3 and 4

Using weights  $[-0.25, 0.55, 0.15]$ :

**Sample 3:**  $x = [1, 0, 0], t = 0$

- $o = -0.25$ , Error = 0.25
- $\Delta w = [0.125, 0, 0]$

**Sample 4:**  $x = [1, 0, 1], t = 1$

- $o = -0.25 + 0 + 0.15 = -0.1$ , Error = 1.1
- $\Delta w = [0.55, 0, 0.55]$

**Accumulated  $\Delta w$ :**  $[0.125 + 0.55, 0, 0 + 0.55] = [0.675, 0, 0.55]$

**Update:**  $[-0.25 + 0.675, 0.55 + 0, 0.15 + 0.55] = [0.425, 0.55, 0.7]$

**Final Answer (Mini-Batch GD, B=2):**

$$w_0 = 0.425, \quad w_1 = 0.55, \quad w_2 = 0.7$$

**Comparison (all after 1 epoch):**

Method	$w_0$	$w_1$	$w_2$	Updates
Batch GD	0.3	0.55	0.55	1
Mini-Batch (B=2)	0.425	0.55	0.7	2
SGD	0.4625	0.15	0.5375	4

## 9 Practical Considerations for SGD

### 9.1 Learning Rate Selection

#### Memory Hook — Remember This!

##### Learning Rate Guidelines:

- **Too large:** Overshoots minimum, may diverge
- **Too small:** Converges very slowly
- **Typical values:** 0.001 to 0.1
- **Start large, decay over time** (learning rate scheduling)




Figure 5: Effect of learning rate on SGD convergence

### 9.2 When SGD Can Escape Local Minima

#### Definition

##### Local Minima Escape:

Because SGD uses different gradients  $\nabla E_d(\mathbf{w})$  for each sample (instead of the true gradient  $\nabla E(\mathbf{w})$ ), the noisy updates can “jump” out of shallow local minima.

This is particularly useful in deep learning where the error surface has many local minima.




Figure 6: SGD’s noisy updates can escape local minima

## 10 Summary: SGD Calculation Checklist

### SGD Step-by-Step Checklist

**For EACH sample in the dataset:**

1. **Forward Pass:** Calculate output

$$o = \sum_{i=0}^n w_i \cdot x_i = w_0 x_0 + w_1 x_1 + \dots + w_n x_n$$

2. **Error Calculation:**

$$e = t - o$$

3. **Weight Updates:** For each weight  $w_i$ :

$$\Delta w_i = \eta \cdot e \cdot x_i$$

$$w_i \leftarrow w_i + \Delta w_i \quad (\text{UPDATE IMMEDIATELY!})$$

4. **Move to next sample** with the **updated weights**

**Repeat for multiple epochs until convergence.**

### Self-Test — Check Your Understanding

#### Quick Practice:

Given:  $\eta = 0.2$ , weights  $[w_0, w_1] = [0.5, -0.3]$ , sample  $x = [1, 2]$ , target  $t = 1$

1. Calculate output  $o$
2. Calculate error  $e$
3. Calculate  $\Delta w_0$  and  $\Delta w_1$
4. What are the updated weights?

#### Answers:

1.  $o = 0.5(1) + (-0.3)(2) = 0.5 - 0.6 = -0.1$
2.  $e = 1 - (-0.1) = 1.1$
3.  $\Delta w_0 = 0.2 \times 1.1 \times 1 = 0.22; \Delta w_1 = 0.2 \times 1.1 \times 2 = 0.44$
4.  $[0.5 + 0.22, -0.3 + 0.44] = [0.72, 0.14]$

## 11 Glossary

Term	Definition
<b>SGD</b>	Stochastic Gradient Descent — updates weights after each sample
<b>Batch GD</b>	Updates weights once after processing all samples
<b>Mini-Batch GD</b>	Updates weights after processing a small batch of samples
<b>Epoch</b>	One complete pass through all training data
<b>Learning Rate (<math>\eta</math>)</b>	Controls the step size of weight updates
<b>Gradient</b>	Direction of steepest increase in error
<b>Local Minimum</b>	A point where error is lower than nearby points but not globally lowest
<b>Convergence</b>	When weights stop changing significantly
<b>Batch Size</b>	Number of samples processed before one weight update