

Regularization Techniques

Parameter Norm Penalties & Dataset Augmentation

Deep Learning for Perception — BSCS, FAST-NUCES

1 Parameter Norm Penalties

Why It Matters

Neural networks have millions of parameters that can easily overfit. Parameter norm penalties add a “cost” to having large weights, forcing models to keep parameters small unless truly important. This is THE most common regularization technique in production deep learning.

1.1 The Core Concept

Analogy

Packing for a trip:

Without penalty: Pack everything “just in case” → huge heavy suitcase, most items unused.

With penalty (\$10/kg): Pack only essentials → light efficient suitcase.

Neural networks: Same idea - make the model “pay” for large weights, so it only uses them when necessary!

1.2 Mathematical Foundation

Definition

Parameter Norm Penalty: Add a penalty term to the loss function that penalizes large parameter values.

General Form:

$$\tilde{L}(\theta) = L(\theta) + \alpha\Omega(\theta)$$

Where:

- $L(\theta)$ = Original loss (prediction error)
- $\Omega(\theta)$ = Regularization term (model complexity)
- α = Regularization strength (hyperparameter)
- $\tilde{L}(\theta)$ = Total loss (what we minimize)

Balance: Model must (1) make good predictions AND (2) keep weights small.

1.3 L2 Regularization (Ridge / Weight Decay)

L2 - Most Common

Idea: Penalize the SQUARED magnitude of weights.

Penalty:

$$\Omega(\theta) = \frac{1}{2} \|w\|_2^2 = \frac{1}{2} \sum_i w_i^2$$

Total Loss:

$$\tilde{L}(w) = L(w) + \frac{\alpha}{2} \sum_i w_i^2$$

Properties:

- Penalizes large weights MORE (squared term)
- Drives weights toward zero (rarely exactly zero)
- Smooth, differentiable
- Also called “Ridge” or “Weight Decay”




Figure 1: L2 penalty grows quadratically with weight magnitude

1.3.1 Weight Decay Interpretation

Key Formula

Gradient of L2-regularized loss:

$$\frac{\partial \tilde{L}}{\partial w} = \frac{\partial L}{\partial w} + \alpha w$$

Gradient descent update:

$$w \leftarrow w - \eta \left(\frac{\partial L}{\partial w} + \alpha w \right)$$

Rearranging:

$$w \leftarrow w(1 - \eta\alpha) - \eta \frac{\partial L}{\partial w}$$

Interpretation:

- $(1 - \eta\alpha)$ = Decay factor (typically 0.999)
- Each update SHRINKS weight by this factor
- Then applies gradient update
- Weight “decays” toward zero!

Typical values:

- $\alpha = 0.0001$ to 0.01
- $\eta = 0.001$ to 0.1
- Decay: $1 - \eta\alpha \approx 0.9999$ to 0.999

Example

Given: $w = 0.8$, $\eta = 0.01$, $\alpha = 0.5$, gradient = 0.3

Solution:

Step 1: Decay factor = $1 - (0.01)(0.5) = 0.995$

Step 2: Apply decay: $0.8 \times 0.995 = 0.796$

Step 3: Apply gradient: $0.796 - 0.01(0.3) = 0.793$

Result: Weight changed from 0.8 → 0.793

- Decay contribution: -0.004
- Gradient contribution: -0.003
- Total: -0.007

1.4 L1 Regularization (Lasso)

L1 - Sparse Solutions

Idea: Penalize ABSOLUTE VALUE of weights.

Penalty:

$$\Omega(\theta) = \|w\|_1 = \sum_i |w_i|$$

Total Loss:

$$\tilde{L}(w) = L(w) + \alpha \sum_i |w_i|$$

Properties:

- Pushes weights to EXACTLY zero (sparse)
- Non-smooth at zero (absolute value)
- Automatic feature selection
- Less common in deep learning




Figure 2: L1 V-shape pushes to zero; L2 U-shape keeps small

1.5 L1 vs L2 Comparison

| Aspect | L1 (Lasso) | L2 (Ridge) |
|-------------------|---------------------------------|-------------------------------|
| Penalty | $\alpha \sum w_i $ | $\frac{\alpha}{2} \sum w_i^2$ |
| Effect | Many $w_i = 0$ exactly | All w_i small |
| Sparsity | YES | NO |
| Feature Selection | Automatic | None |
| Best For | Many irrelevant features | All features contribute |
| Example | Text (10K words, 100 important) | Images (all pixels) |
| Popularity | Less common | MOST COMMON |

Quick Reference

Quick Guide:

L2 (Default):

- Penalty: $\frac{\alpha}{2} \sum w_i^2$
- Update: $w \leftarrow w(1 - \eta\alpha) - \eta\nabla L$
- Use 90% of the time
- Typical α : 0.01 or 0.001

L1 (Special Cases):

- Penalty: $\alpha \sum |w_i|$
- Creates sparse models
- Use when want feature selection
- Typical α : 0.001

Implementation (PyTorch):

```
optimizer = torch.optim.Adam(
    model.parameters(),
    lr=0.001,
    weight_decay=0.01  # L2 alpha
)
```

2 Dataset Augmentation

Why It Matters

MORE DATA = BETTER MODELS. But collecting real data is expensive. Dataset augmentation creates NEW training examples by transforming existing data in label-preserving ways. Can effectively 10x-100x your dataset FOR FREE! Used in every state-of-the-art vision model.

2.1 The Core Idea

Analogy

Learning to recognize cars:

Without augmentation: 100 photos, all front view, daylight → fails on sideways or night photos.

With augmentation: Same 100 photos rotated, brightness adjusted, flipped → sees cars from all angles/lighting → recognizes ANY car!

Magic: Created 1000 training examples from 100 originals!

Definition

Dataset Augmentation: Generate new training examples by applying label-preserving transformations.

Key Principles:

1. **Preserves label:** Rotated cat is still a cat
2. **Increases diversity:** Different views/conditions
3. **During training only:** Each epoch sees new versions
4. **No test augmentation:** Use originals for evaluation

Benefits:

- Reduces overfitting
- Improves generalization
- Effectively increases dataset size
- Robust to real-world variations
- FREE!

2.2 Image Augmentation Techniques

2.2.1 Geometric Transformations

| Transform | Use When | Avoid When |
|---------------------------------------|-------------------|-----------------------|
| Rotation ($\pm 5\text{--}45^\circ$) | Any angle OK | Text, street signs |
| Horizontal Flip | Symmetric objects | Digits (6 becomes 9!) |
| Vertical Flip | Satellite images | Natural images |
| Translation | Object anywhere | Centered only |
| Zoom (80–120%) | Varying distances | Fixed-size objects |
| Random Crop | Large images | Small images |




Figure 3: Geometric augmentations preserve labels

2.2.2 Photometric Transformations

| Transform | Effect | Typical Range |
|----------------|--------------------|--------------------------|
| Brightness | Lighter/darker | $\times [0.8, 1.2]$ |
| Contrast | More/less contrast | $\times [0.7, 1.3]$ |
| Saturation | Color intensity | $\times [0.5, 1.5]$ |
| Hue Shift | Color rotation | $[-10^\circ, +10^\circ]$ |
| Gaussian Noise | Random pixels | $\sigma = 0.01$ |
| Blur | Smoothing | kernel 3×3 |

2.2.3 Advanced Techniques

Modern Augmentations

1. CutOut / Random Erasing

- Randomly mask rectangular regions
- Forces use of all image parts

2. MixUp

- Blend two images: $x_{new} = \lambda x_1 + (1 - \lambda)x_2$
- Label: $y_{new} = \lambda y_1 + (1 - \lambda)y_2$

3. CutMix

- Cut region from image 1, paste into image 2
- Label proportional to areas

4. AutoAugment

- Automatically learn best policy
- State-of-the-art results

2.3 Domain-Specific Guidelines

When to Use What

Medical Imaging:

- YES: Rotation, scaling, elastic deformation
- NO: Color changes (diagnostic!), flips ($\text{left} \neq \text{right}$)

Text/Documents:

- YES: Small rotation ($\pm 2^\circ$), brightness, perspective
- NO: Horizontal flip (unreadable!), large rotation

Natural Images:

- YES: Rotation ($\pm 15^\circ$), h-flip, crop, color jitter
- NO: Vertical flip (usually)

2.4 Implementation

PyTorch Implementation

```
from torchvision import transforms

# Training augmentation
train_transform = transforms.Compose([
    transforms.RandomRotation(15),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomResizedCrop(224, scale=(0.8, 1.0)),
    transforms.ColorJitter(brightness=0.2,
                          contrast=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                      std=[0.229, 0.224, 0.225])
])

# NO augmentation for validation/test
val_transform = transforms.Compose([
    transforms.Resize(256),
```

```
transforms.CenterCrop(224),  
transforms.ToTensor(),  
transforms.Normalize(mean=[0.485, 0.456, 0.406],  
                    std=[0.229, 0.224, 0.225])  
])
```

2.5 Best Practices

Quick Reference

Golden Rules:

DO:

- Augment TRAINING only (not test!)
- Apply multiple together
- Start conservative, increase gradually
- Visualize samples (check they make sense)
- Combine with dropout and L2

DON'T:

- Augment test data (biases evaluation!)
- Change the label
- Over-augment (unrecognizable)
- Use same every epoch

Safety Check: “Would a human still label this correctly?”

- YES → Safe
- NO → Don't use

Augmentation Strength:

- Light: Rotation $\pm 5^\circ$, brightness $\pm 10\%$
- Medium: Rotation $\pm 15^\circ$, brightness $\pm 20\%$, flip
- Heavy: Rotation $\pm 30^\circ$, crop, flip, color, cutout

Start with Medium, adjust if:

- Still overfitting → increase
- Training accuracy too low → decrease

Impact Example

Scenario: Cat vs Dog classifier, 1000 images

| Setup | Train | Val | Gap |
|---------------------------|-------|-----|-----|
| No augmentation | 95% | 72% | 23% |
| Light (flip) | 92% | 78% | 14% |
| Medium (flip+rotate+crop) | 88% | 83% | 5% |
| Heavy (all) | 85% | 84% | 1% |

Best: Medium augmentation

- +11% validation ($72 \rightarrow 83\%$)
- 78% gap reduction ($23 \rightarrow 5\%$)
- Equivalent to 5-10x more data!

3 Summary

Key Takeaways

Parameter Norm Penalties

L2 (Weight Decay) - DEFAULT:

- Penalty: $\frac{\alpha}{2} \sum w_i^2$
- Update: $w \leftarrow w(1 - \eta\alpha) - \eta\nabla L$
- All weights \rightarrow small
- Use 90% of the time
- Typical α : 0.01 or 0.001

L1 (Lasso) - SPARSE:

- Penalty: $\alpha \sum |w_i|$
- Many weights \rightarrow exactly 0
- Feature selection
- Typical α : 0.001

Dataset Augmentation

Common Techniques:

- Geometric: Rotation, flip, crop, zoom
- Photometric: Brightness, contrast, color
- Advanced: CutOut, MixUp, CutMix

Rules:

- Training only (not test!)
- Preserve labels
- Human should still recognize
- Combine with L2 + dropout

Impact:

- 5-15% accuracy improvement
- Dramatically reduces overfitting
- Equivalent to 5-10x more data
- FREE!

Standard Recipe

```
Data Augmentation (flip+rotate+crop+color)
+ L2 Regularization (alpha=0.01)
+ Dropout (p=0.5)
+ Early Stopping
= Robust Model
```

End of Notes

Deep Learning for Perception — FAST-NUCES